

# Memory Bound Computing

**Francesc Alted**

Freelance Consultant

<http://www.blosc.org/professional-services.html>

Advanced Scientific Programming in Python

Reading, UK

September, 2016

“No sensible decision can be made any longer without taking into account not only the world as it is, but the world as it will be.”

— *Isaac Asimov*

“No sensible decision can be made any longer without taking into account not only the **computer** as it is, but the **computer** as it will be.”

— My own rephrasing

# Overview


- The Starving CPU problem
- Recent trends in computer architecture
- Handling Big Data: storing and processing as much data as possible with your existing resources

*“Across the industry, today’s chips are largely able to execute code faster than we can feed them with instructions and data.”*

– Richard Sites, after his article  
“It’s The Memory, Stupid!”,  
Microprocessor Report, 10(10), 1996


# The Starving CPU Problem

# Memory Access Time vs CPU Cycle Time



The gap is wide and still opening!

Book in  
2009



# The Status of CPU Starvation in 2016

- Memory latency is much slower (between 250x and 1000x) than processors.
- Memory bandwidth is improving at a better rate than memory latency, but it is also slower than processors (between 30x and 100x).

# CPU Caches to the Rescue

- CPU cache latency and throughput are much better than memory
- However: the faster they run the smaller they must be (because of heat dissipation problems)




# Computer Architecture Evolution

Up to end 80's

90's and 2000's

2010's



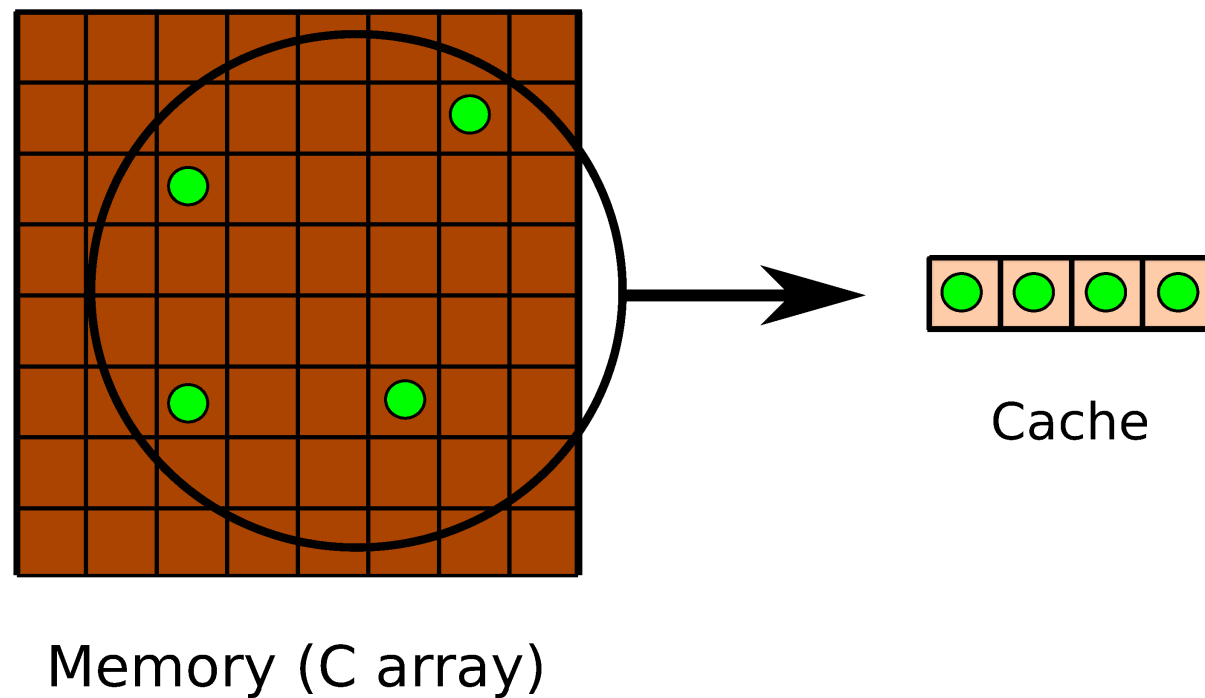
# When CPU Caches Are Effective?

Mainly in a couple of scenarios:

- Time locality: when the dataset is reused
- Spatial locality: when the dataset is accessed sequentially


# Time Locality

Parts of the dataset are reused



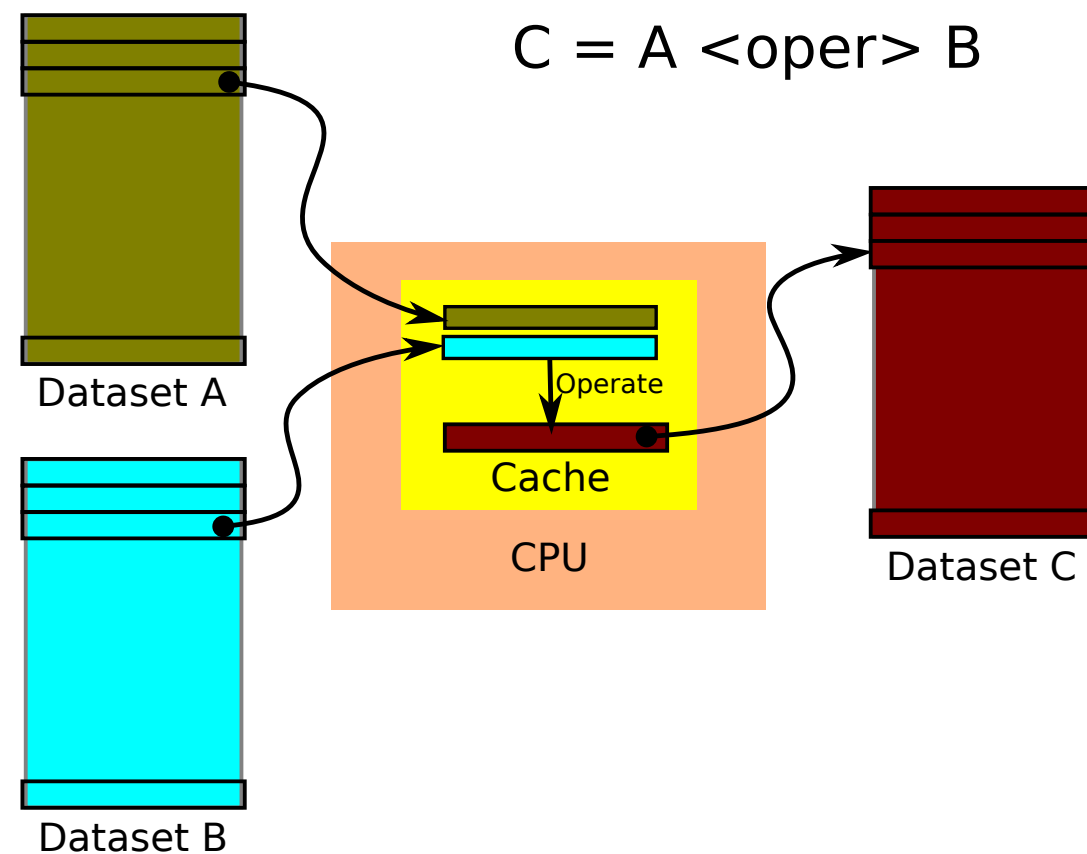
# Spatial Locality

Dataset is accessed sequentially



# The Blocking Technique

When accessing disk or memory, get a **contiguous** block that fits in CPU cache, operate upon it and **reuse** it as much as possible.



Use this extensively to leverage **spatial** and **temporal** localities


*“Across the industry, today’s chips are largely able to execute code faster than we can feed them with instructions and data. There are no longer performance bottlenecks in the floating-point multiplier or in having only a single integer unit. The real design action is in memory subsystems— caches, buses, bandwidth, and latency.”*

*“Over the coming decade, memory subsystem design will be the only important design issue for microprocessors.”*

— Richard Sites, after his article “It’s The Memory, Stupid!”,  
Microprocessor Report, 10(10), 1996

# Trends in Computer Architecture


# We Are In A Multicore Age



- This requires special programming measures to leverage all its potential: threads, multiprocessing




# SIMD: Single Instruction, Multiple Data




More operations in the same CPU clock

# Forthcoming Trends in CPU


## More Cores



## Wider Vectors




## CPU+GPU Integration




# Hierarchy of Memory

## By 2018 (Educated Guess)



# Efficient Data Containers for In-Memory Storage

# Which Structure Is More Efficient for Computing?




Attr: Jake Vanderplas

# numexpr

- numexpr is a computational engine for NumPy that makes a sensible use of the memory hierarchy for better performance
- Other libraries normally use numexpr when it is time to accelerate large computations
- PyTables, pandas and bcolz (among others) can all leverage numexpr


# Time To Answer Pending Questions

Time to evaluate polynomial (1 thread)



# Computing with NumPy

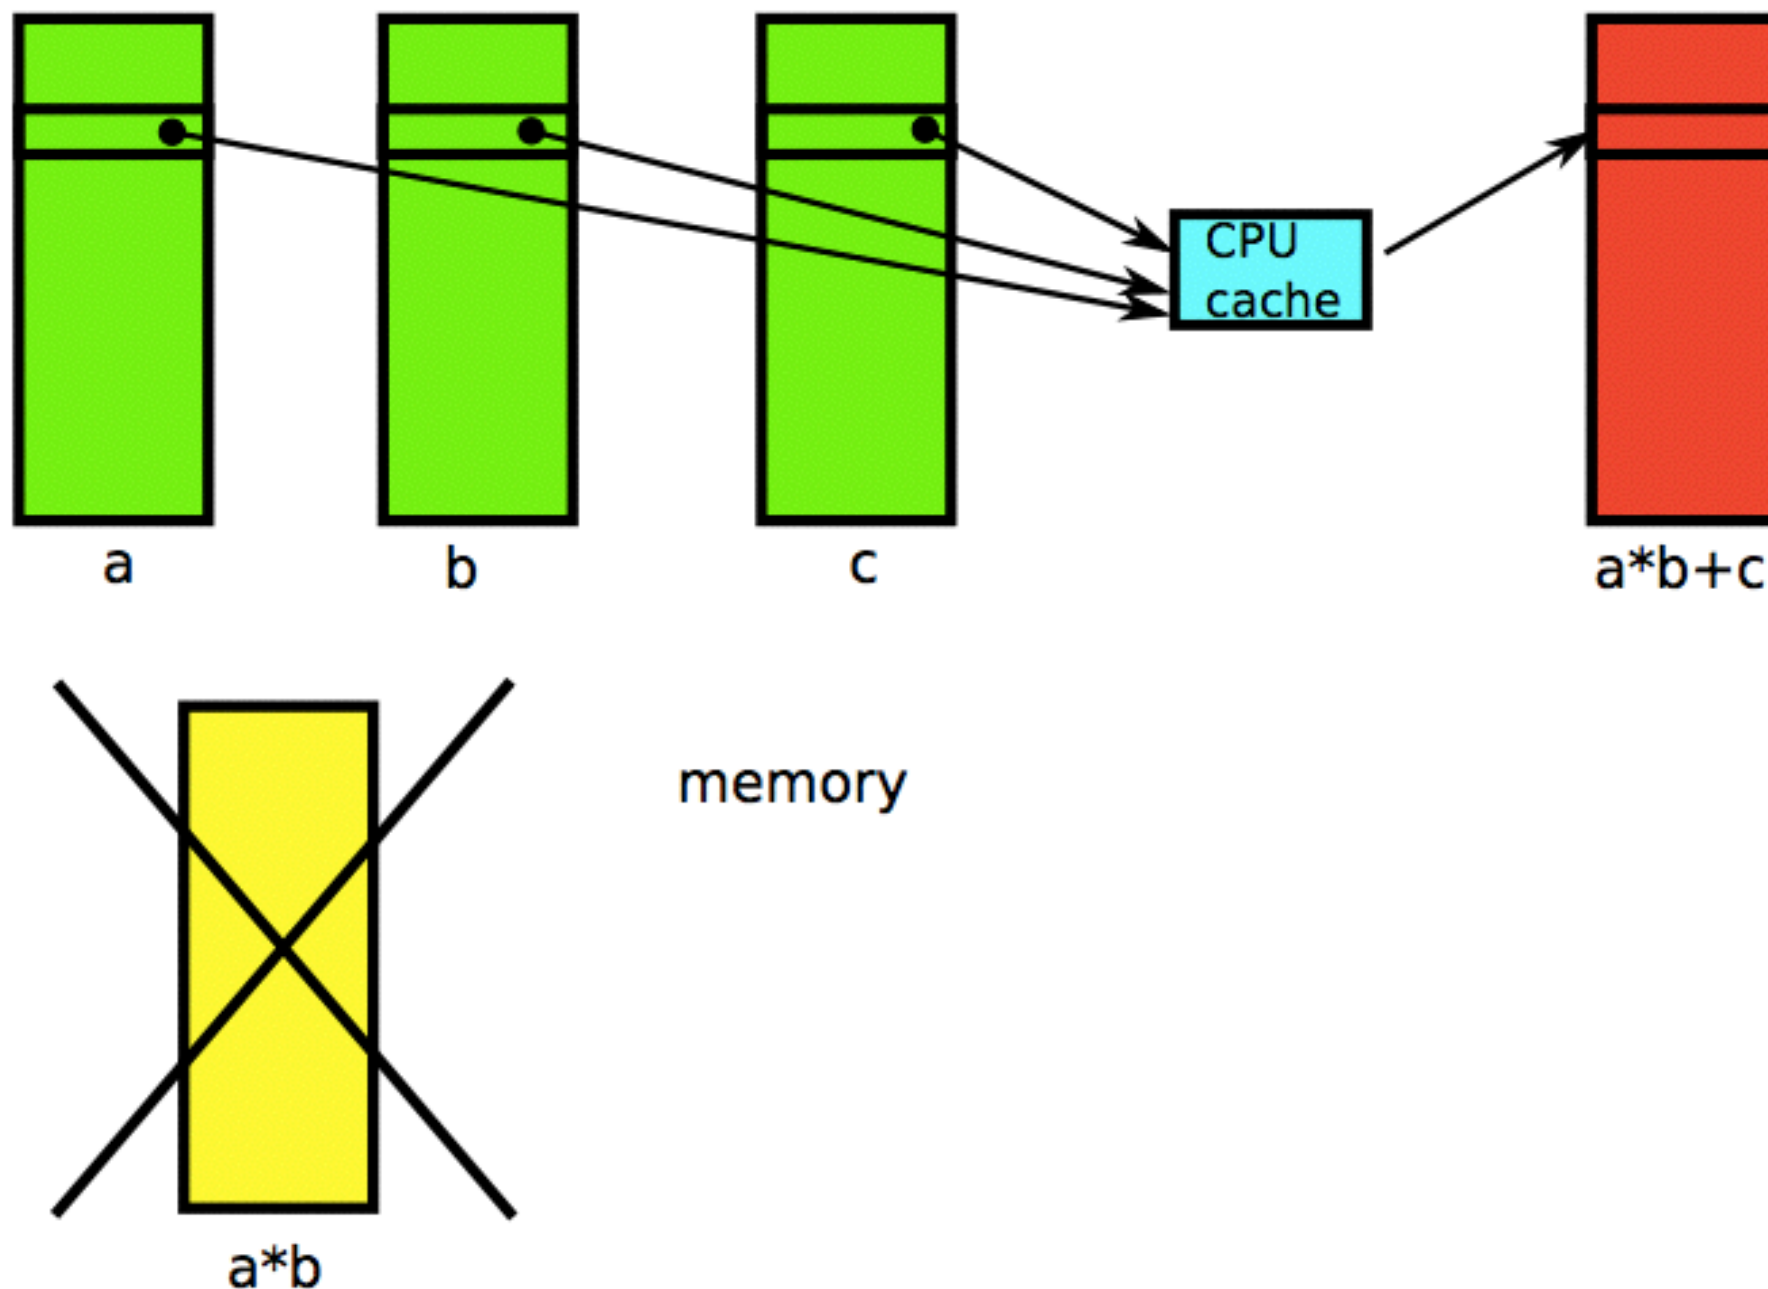
Temporaries go to memory





# Computing with numexpr


Temporaries stay in cache



# Multithreaded numexpr and Beyond: Numba

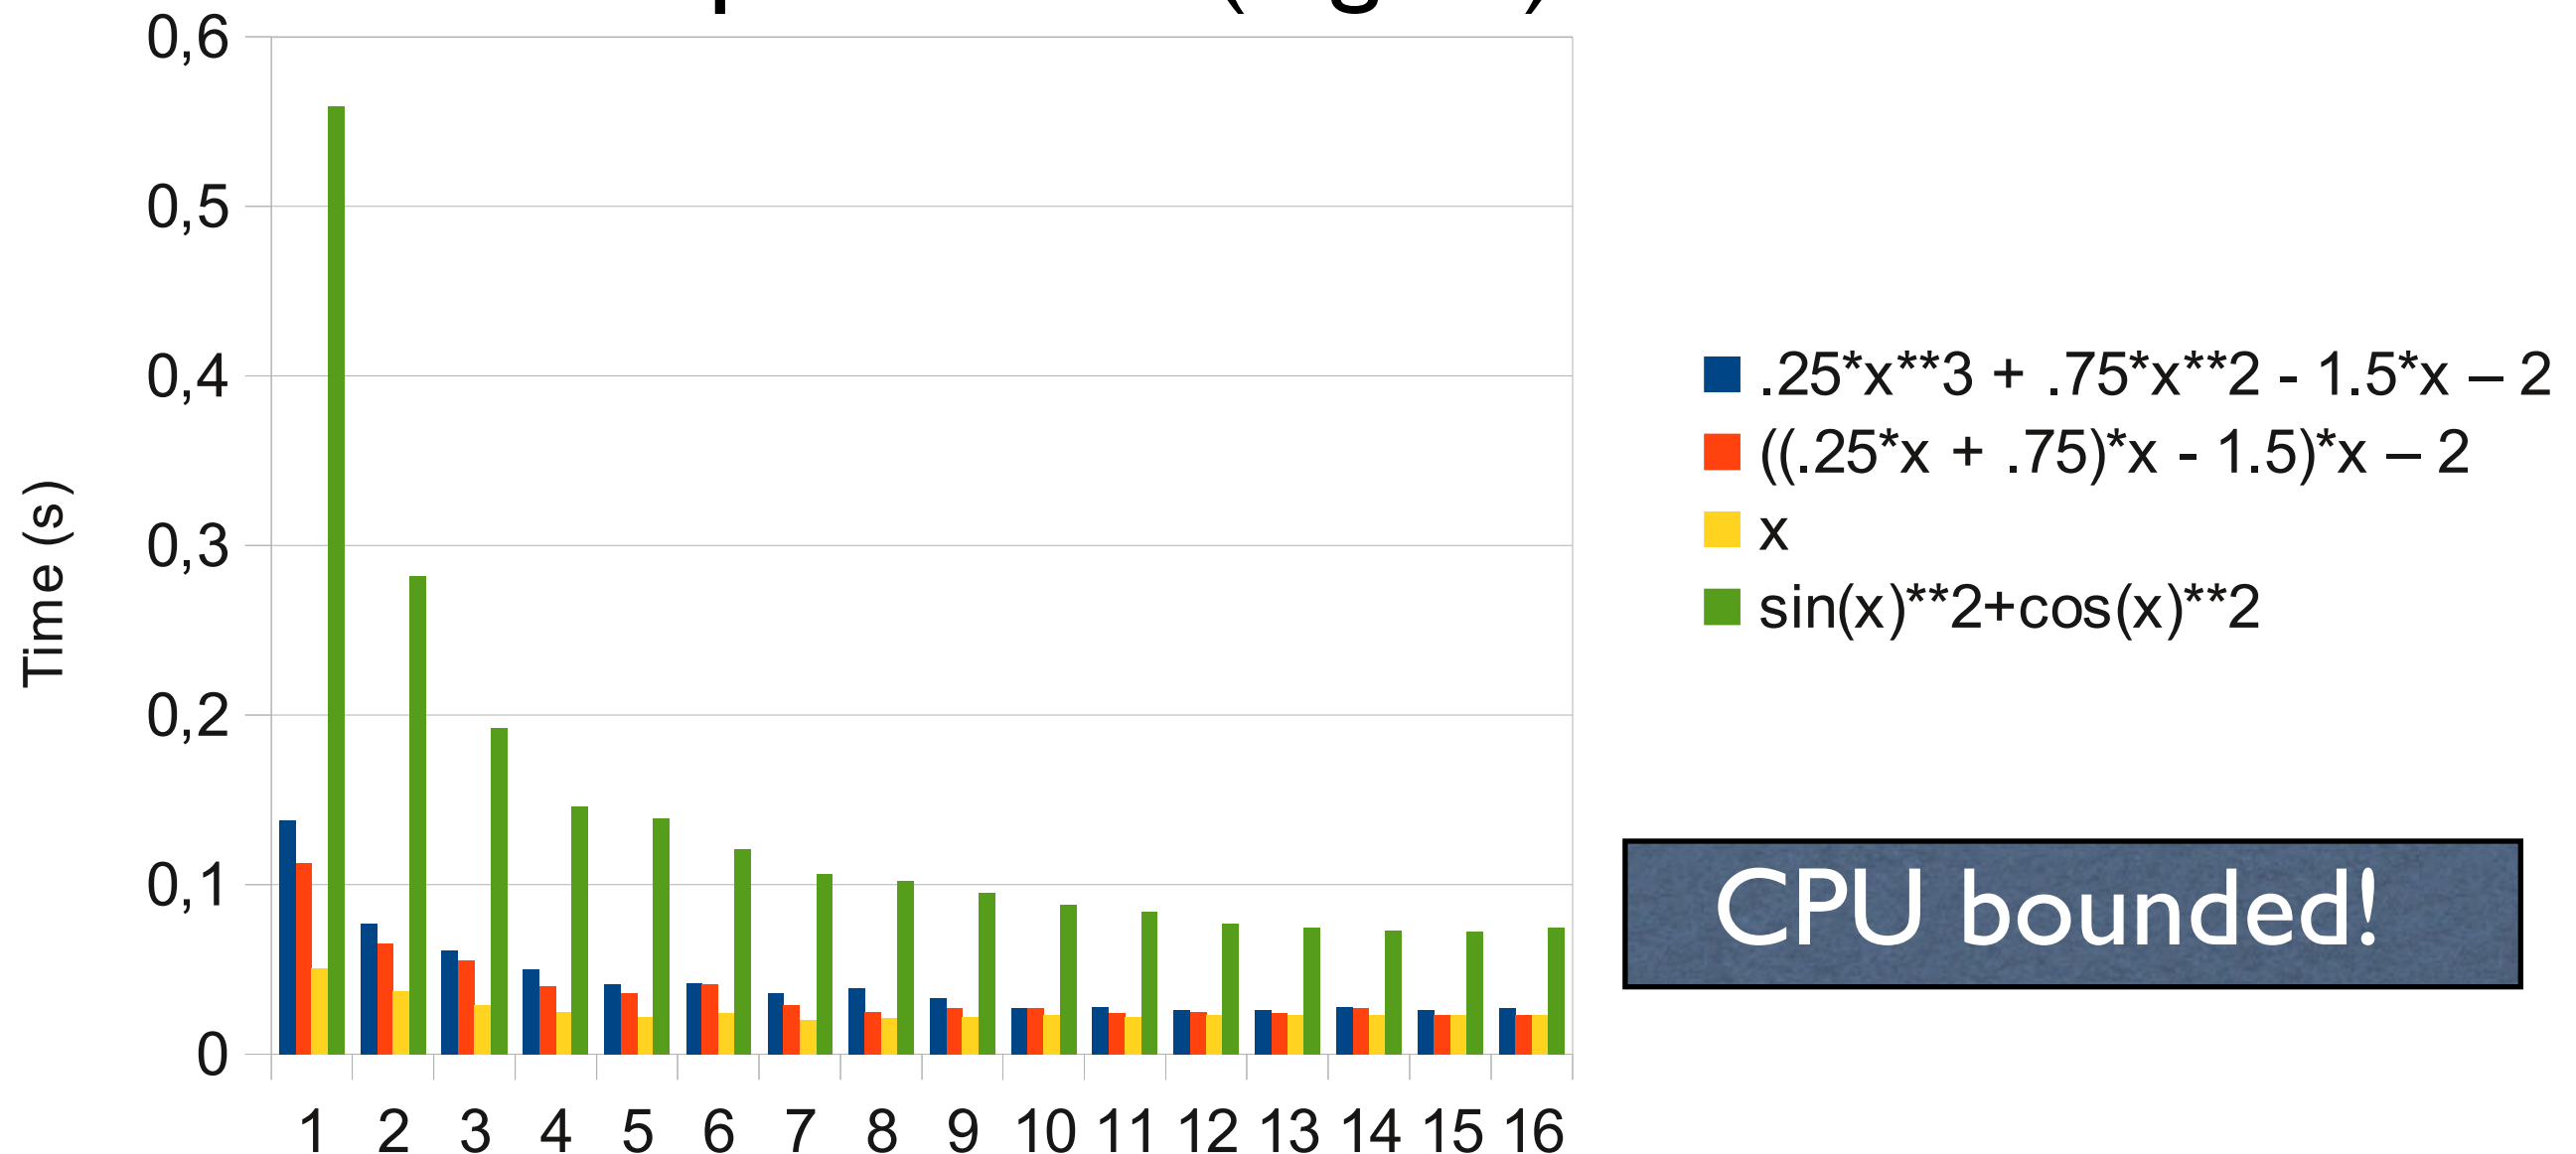
# numexpr Allows Multithreading for Free

numexpr with 16 (logical) cores



# Transcendental Functions

numexpr with 16 (logical) cores



CPU bounded!

# Numexpr Limitations

- Numexpr only implements element-wise operations, i.e. '**a\*b**' is evaluated as:

```
for i in range(N):
```

```
    c[i] = a[i] * b[i]
```

- In particular, it cannot deal with things like:

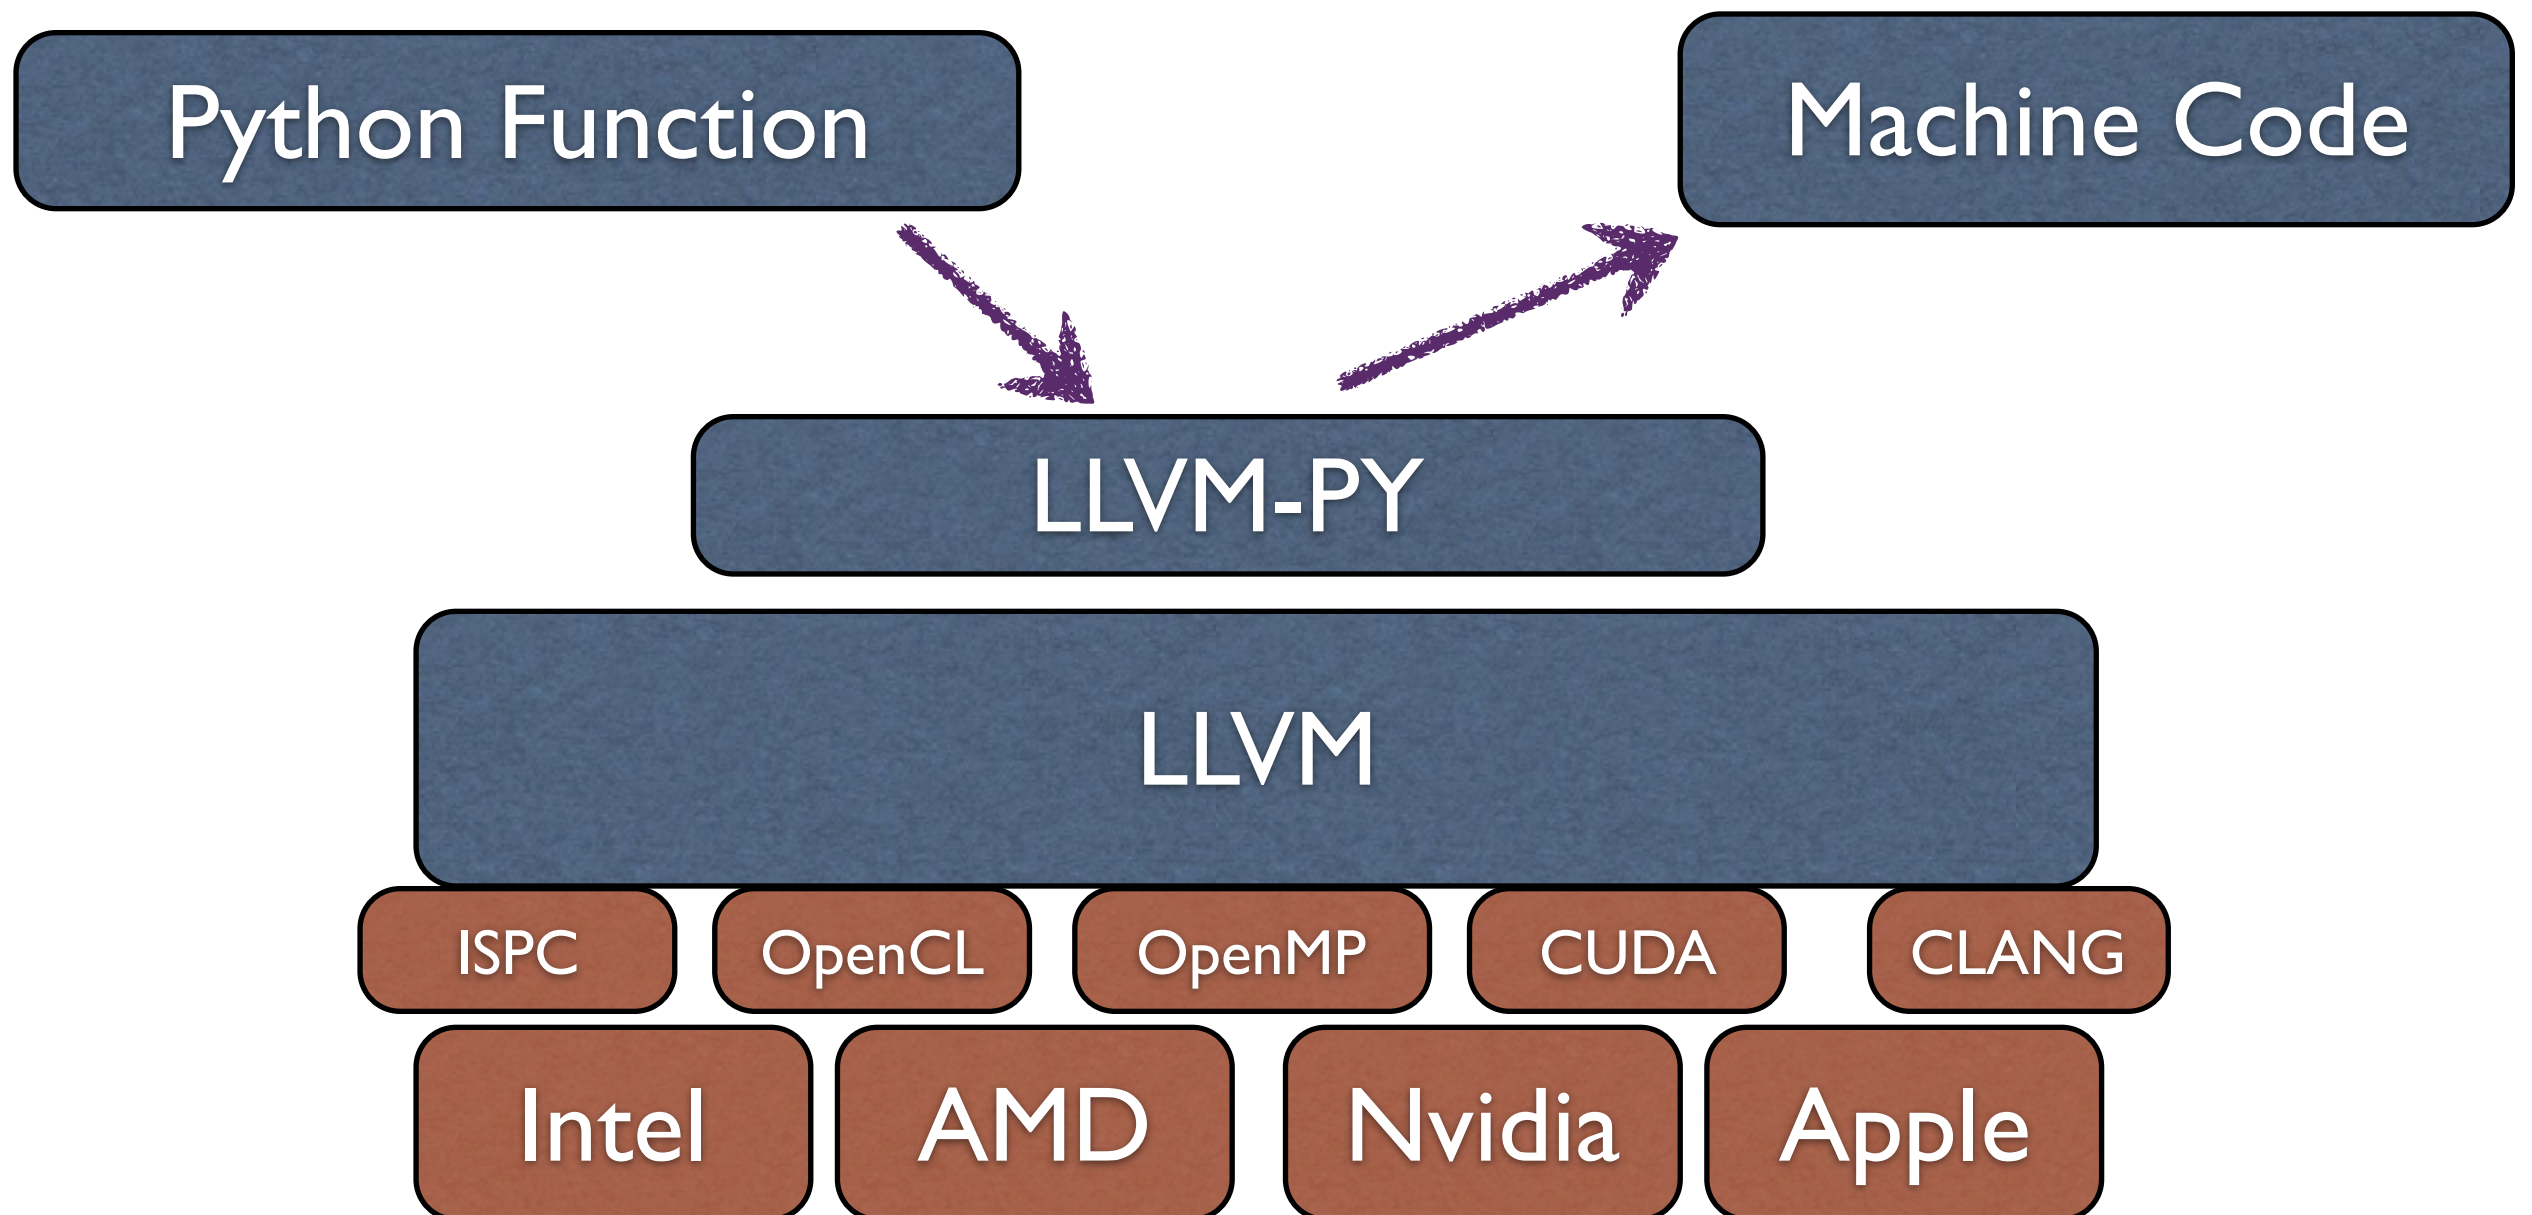
```
for i in range(N):
```

```
    c[i] = a[i-1] + a[i] * b[i]
```

# Numba: Overcoming numexpr Limitations

- Numba is a JIT that can translate a subset of the Python language into machine code
- It uses LLVM infrastructure behind the scenes
- Can achieve similar or better performance than numexpr, but with more flexibility

# How Numba works




# Trends in Computer Storage




The growing gap between DRAM and HDD is facilitating the introduction of new SSD devices

## The DRAM/HDD Speed Gap



From: *Solid State Drives in the Enterprise*  
by Objective Analysis



PCIe SSD



M.2 SSD



BGA SSD

# Latency Numbers Every Programmer Should Know


## Latency Comparison Numbers

-----					
L1 cache reference	0.5	ns			
Branch mispredict	5	ns			
L2 cache reference	7	ns			14x L1 cache
Mutex lock/unlock	25	ns			
Main memory reference	100	ns			20x L2 cache, 200x L1 cache
Read 4K randomly from memory	1,000	ns	0.001	ms	
Compress 1K bytes with Zippy	3,000	ns			
Send 1K bytes over 1 Gbps network	10,000	ns	0.01	ms	
Read 4K randomly from SSD*	150,000	ns	0.15	ms	
Read 1 MB sequentially from memory	250,000	ns	0.25	ms	
Round trip within same datacenter	500,000	ns	0.5	ms	
Read 1 MB sequentially from SSD*	1,000,000	ns	1	ms	4X memory
Disk seek	10,000,000	ns	10	ms	20x datacenter roundtrip
Read 1 MB sequentially from disk	20,000,000	ns	20	ms	80x memory, 20X SSD
Send packet CA->Netherlands->CA	150,000,000	ns	150	ms	

Source: Jeff Dean and Peter Norvig (Google), with some additions

[http://www.eecs.berkeley.edu/~rcs/research/interactive\\_latency.html](http://www.eecs.berkeley.edu/~rcs/research/interactive_latency.html)

# Reference Time vs Transmission Time



$t_{ref} \sim t_{trans} \Rightarrow$  optimizes memory access

# Not All Storage Layers Are Created Equal

**Memory:**  $t_{\text{ref}}$ : 100 ns /  $t_{\text{trans}}$  (**1 KB**): ~100 ns

**Solid State Disk:**  $t_{\text{ref}}$ : 10  $\mu$ s /  $t_{\text{trans}}$  (**4 KB**): ~10  $\mu$ s

**Mechanical Disk:**  $t_{\text{ref}}$ : 10 ms /  $t_{\text{trans}}$  (**1 MB**): ~10 ms

The slower the media, the larger the block that is worth to transmit

But essentially, a blocked data access is mandatory for speed!

# We Need More Data Blocking In Our Infrastructure!

- Not many data containers leveraging the blocking technique exist yet, but a handful (e.g. HDF5, bcolz or zarr) do
- With blocked access we can use persistent media (disk) as it is ephemeral (memory) and the other way around -> independency of media!
- **No silver bullet:** we won't be able to find a single container that makes everybody happy; it's all about tradeoffs

Can We Get Better  
Bandwidth Than  
Hardware Allows?



# Compression for Random & Sequential Access in SSDs

Performance Specification	Incompressible Data	Compressible Data
Sequential Write Bandwidth (Mbps)	235	520
Sequential Read Bandwidth (Mbps)	550	550
Random Write (IOPS)	16,500 (65MB/s)	60,000 (240MB/s)
Random Read (IOPS)	46,000 (180MB/s)	50,000 (200MB/s)

3. Source: Intel® Solid-State Drive 520 Series Product Specification; Random reads based on 4KB Queue Depth 32

- Compression does help performance!

# Compression for Random & Sequential Access in SSDs

Performance Specification	Incompressible Data	Compressible Data
Sequential Write Bandwidth (Mbps)	235	520
Sequential Read Bandwidth (Mbps)	550	550
Random Write (IOPS)	16,500 (65MB/s)	60,000 (240MB/s)
Random Read (IOPS)	46,000 (180MB/s)	50,000 (200MB/s)


3. Source: Intel® Solid-State Drive 520 Series Product Specification; Random reads based on 4KB Queue Depth 32

- Compression does help performance!
- However, limited by SATA bandwidth




# Leveraging Compression Straight To CPU

Less data needs to be transmitted to the CPU




Transmission + decompression faster than direct transfer?



When we have a fast enough compressor we can get rid of the limitations of the bus bandwidth.


**How to get maximum compression performance?**

# NETCDF C LIB: Write TP vs File Size cksize\_T:23 cksize\_YX:128



Thanks to: Rui Yang, Pablo Larraondo (NCI Australia)

Example with actual data (satellite images):  
Blosc compression does not degrade I/O performance



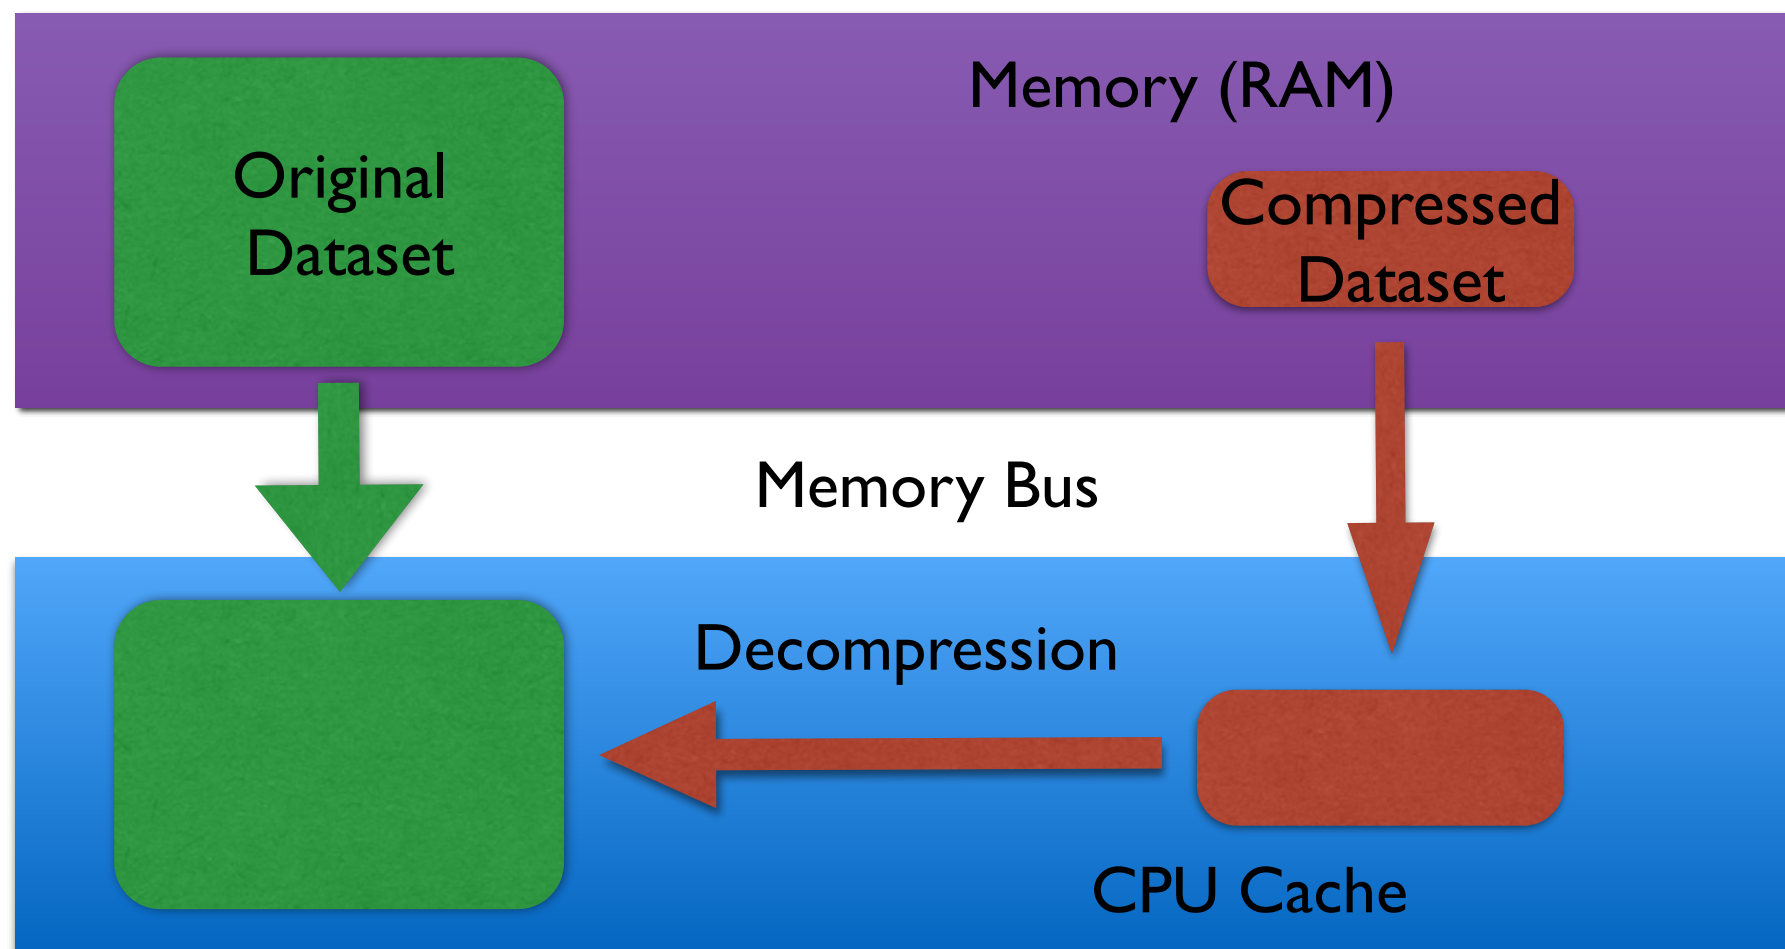
Thanks to: Rui Yang, Pablo Larraondo (NCI Australia)

Reading satellite images:  
Blosc decompression accelerates I/O

Can CPU-based  
Compression Alleviate  
The Memory Bottleneck?

# Improving RAM Speed?

Less data needs to be transmitted to the CPU



Transmission + decompression faster than direct transfer?

# We can try, but certain conditions must be met


- We must follow the principles we have seen:
  - Data Containers leveraging the blocking technique (better cache usage)
  - Using (extremely fast) compression per every block

# Principles of Blosc

- Split data chunks in blocks internally (better cache utilization)
- Supports Shuffle and BitShuffle filters
- Uses parallelism at two levels:
  - Use multicores (multithreading)
  - Use SIMD in Intel/AMD processors (SSE2, AVX2) and ARM (NEON)




# The Shuffle filter



- Shuffle works at byte level, and works well for integers or floats that vary smoothly
- There is also support for a BitShuffle filter that works at bit level

# Blosc: Compressing Faster Than *memcpy()*




# **bcolz:** a Data Container that Leverages the Blocking Technique

NumPy container

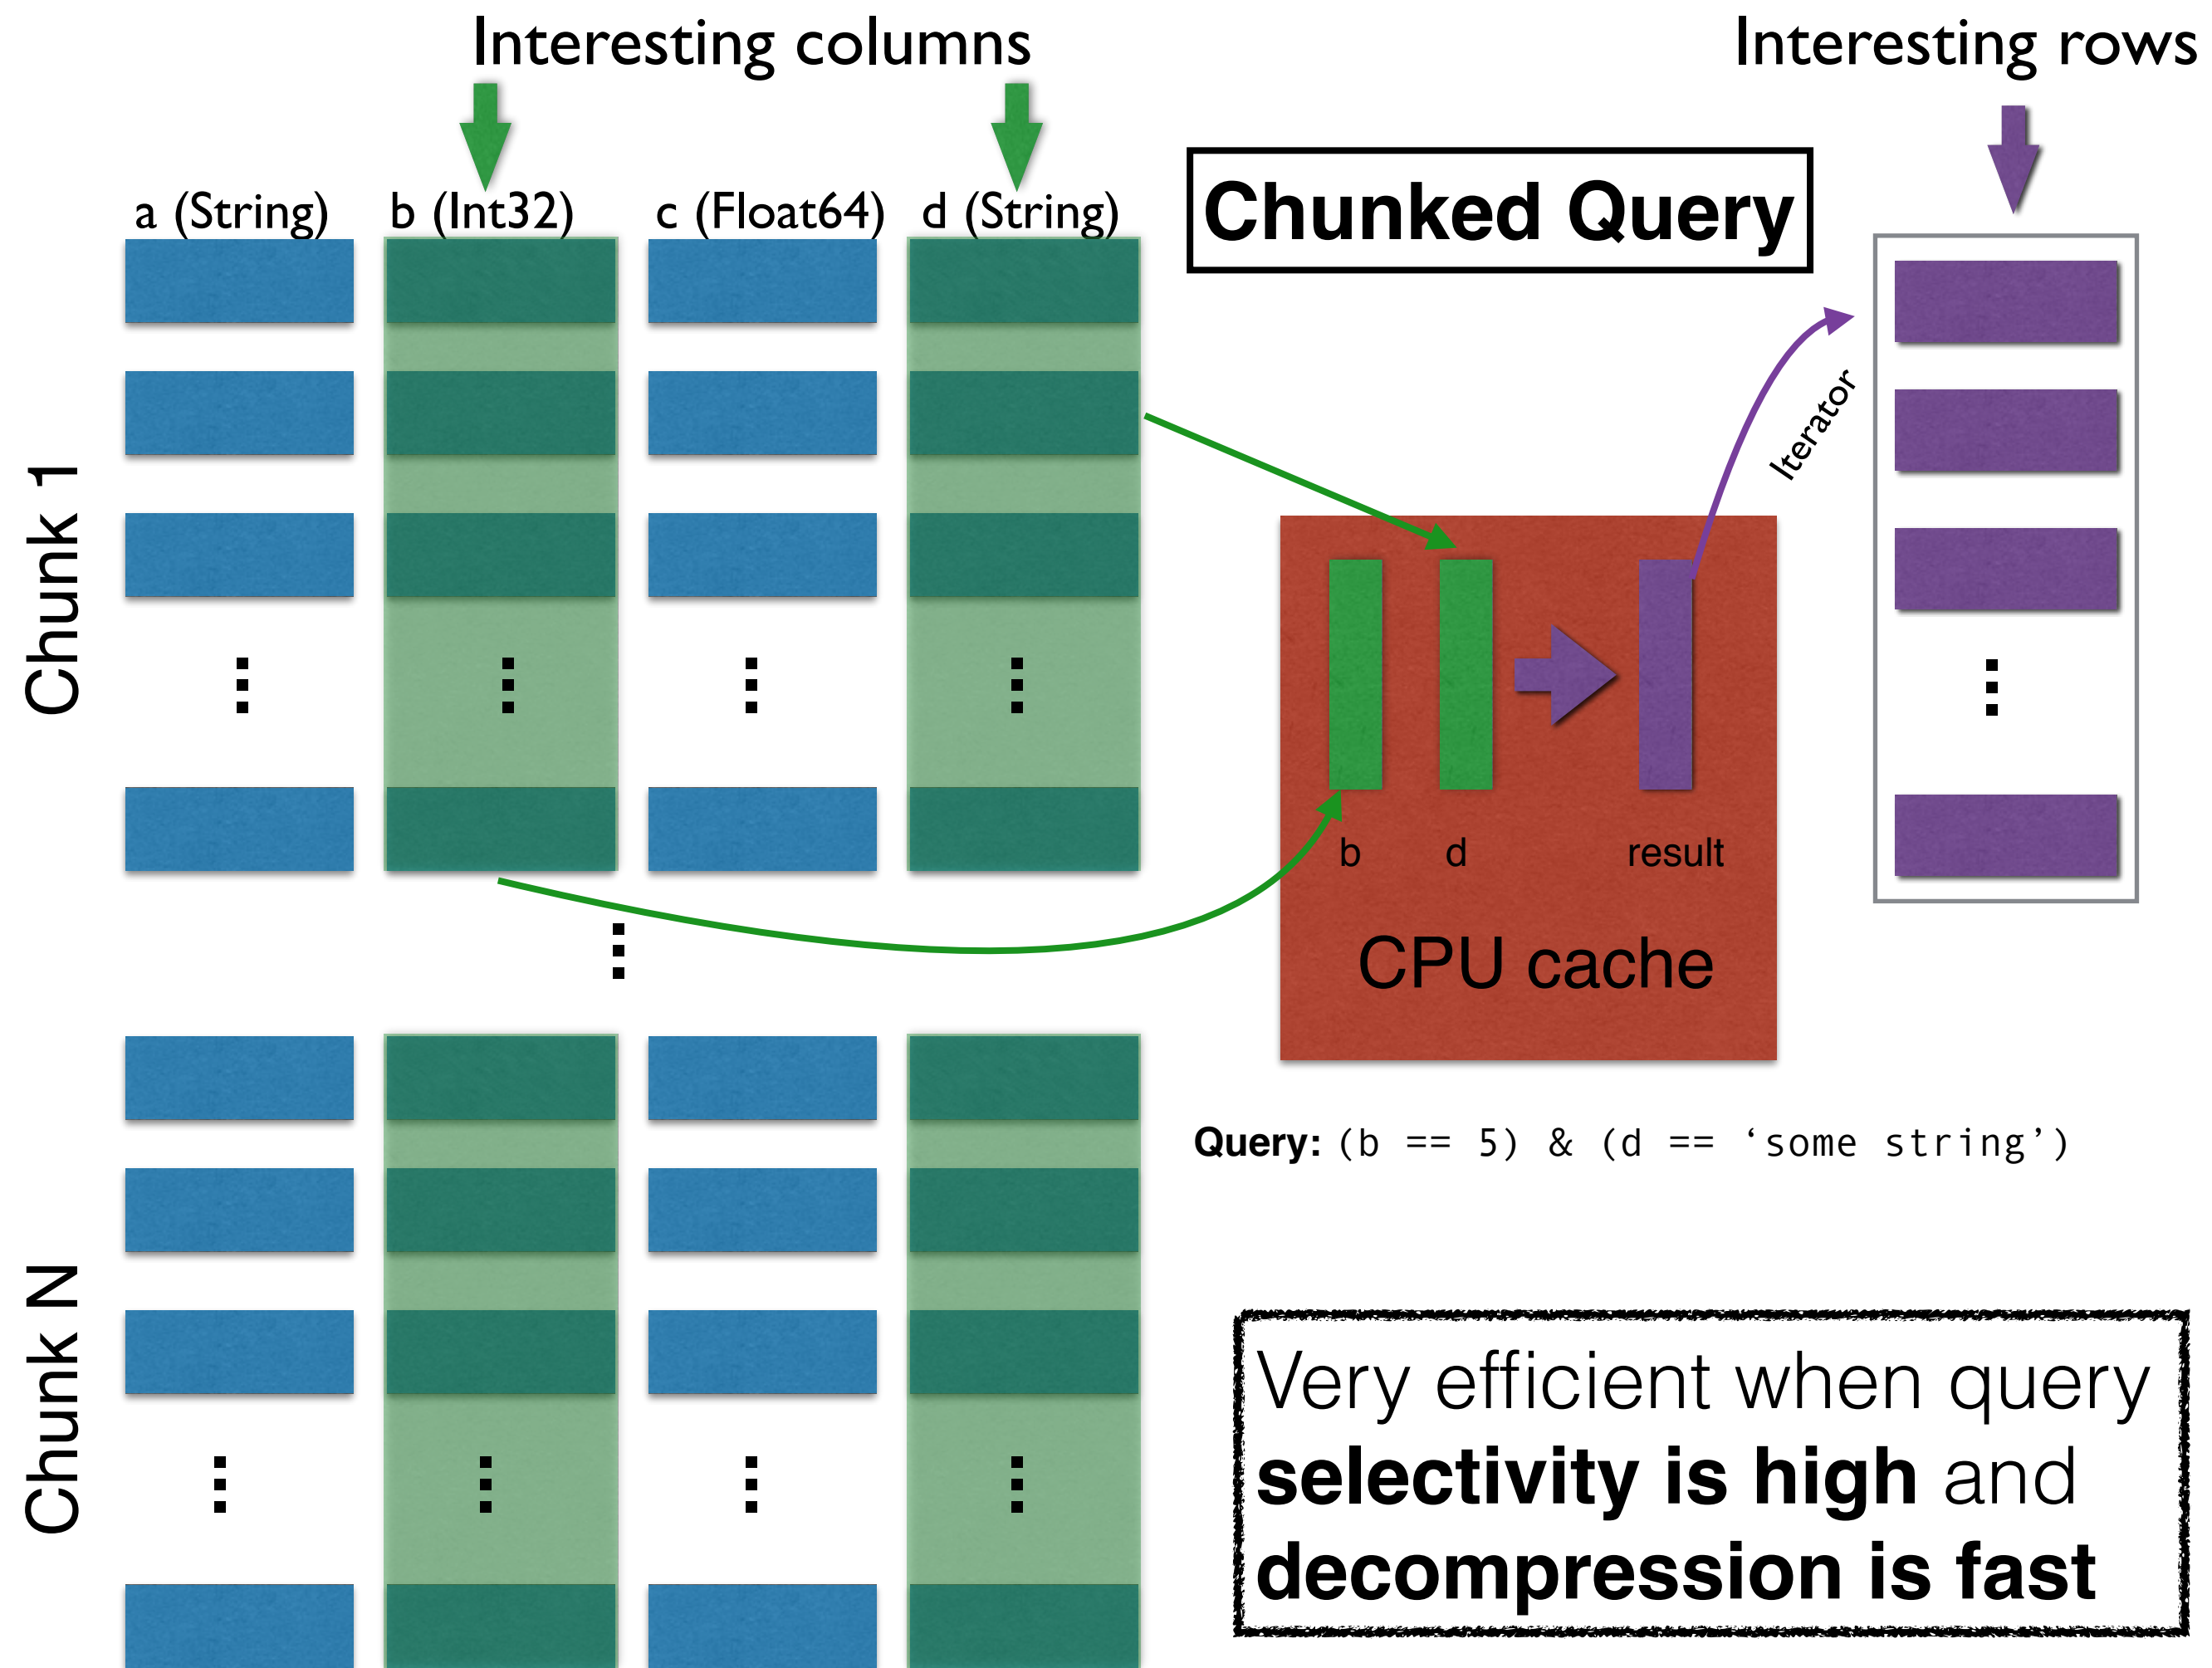


Contiguous memory

bcolz container




Discontiguous memory



# Query Times in bcolz

Recent server (Intel Xeon Skylake, 4 cores)  
Compression **speeds** things up




Reference: <https://github.com/Blosc/movielens-bench/blob/master/querying-ep14.ipynb>

# Query Times in bcolz


4-year old laptop (Intel Ivy-Bridge, 2 cores)

Compression still **speeds** things up




# Query Times in bcolz

2010 laptop (Intel Core2, 2 cores)  
Compression **slows** things down



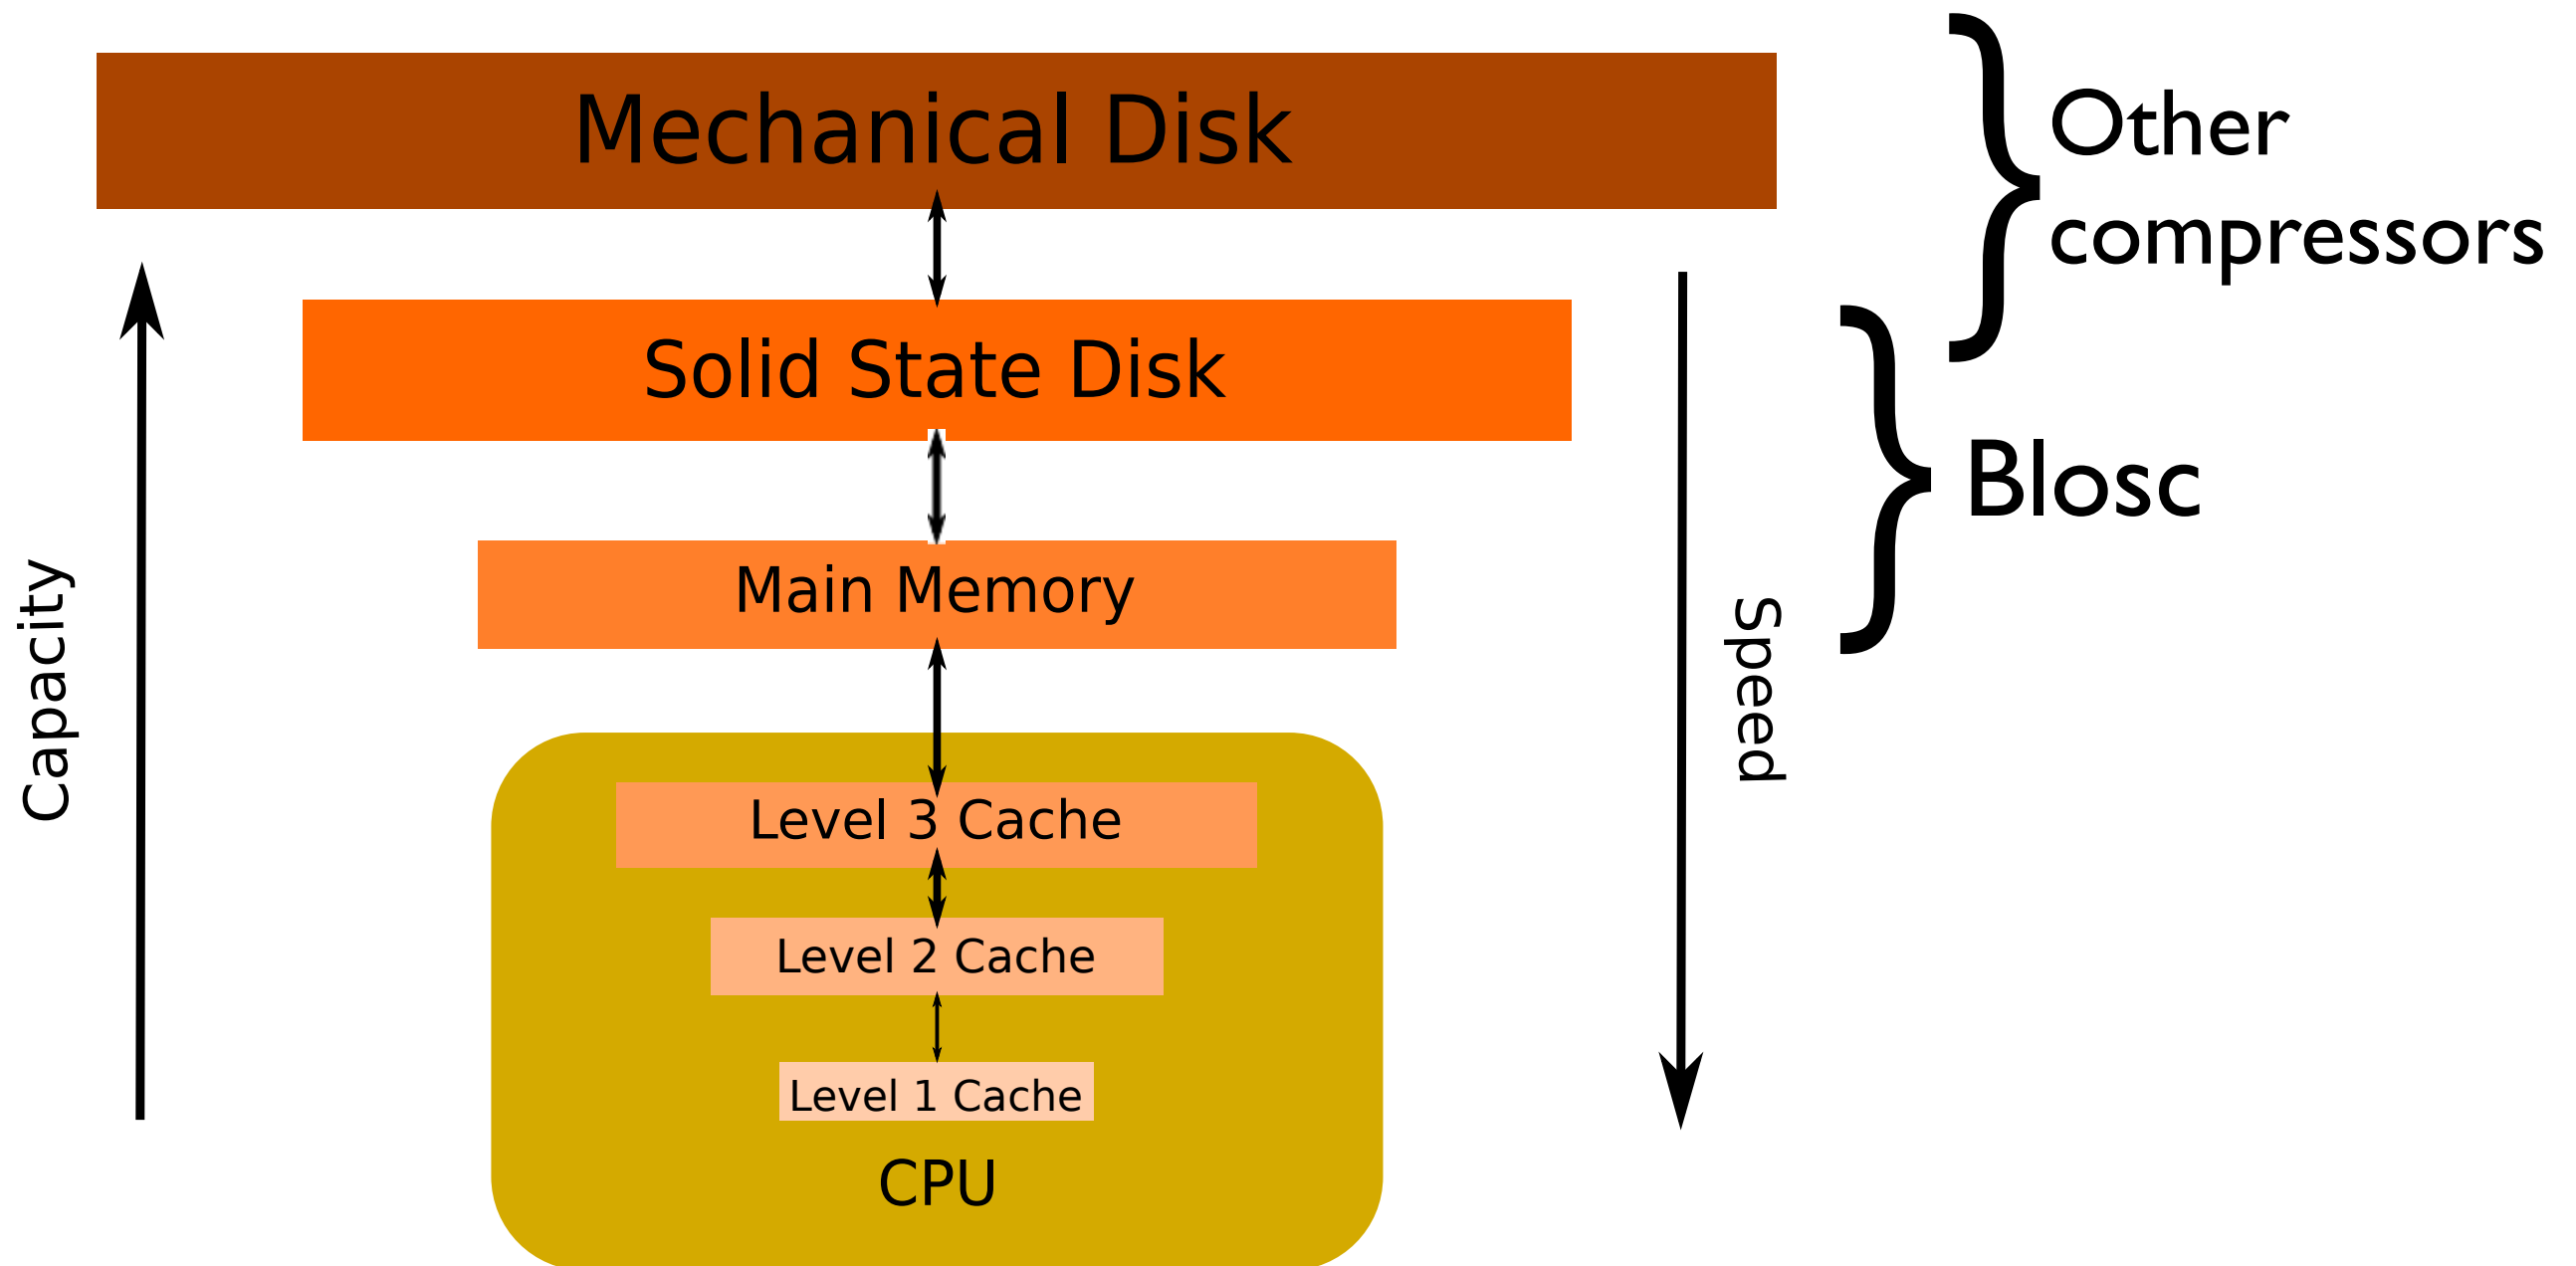
# Sizes in bcolz



Do not forget compression main strength:  
**We can store more data while using same resources**



# Accelerating I/O With Blosc



“Blosc compressors are the fastest ones out there at this point; there is no better publicly available option that I'm aware of. That's not just ‘yet another compressor library’ case.”

— *Ivan Smirnov*  
(*advocating for Blosc inclusion in h5py*)

# Compression matters!

# Closing Notes

- Due to the evolution in computer architecture, compression can be effective for two reasons:
  - We can work with more data using the same resources.
  - We can reduce the overhead of compression to near zero, and even beyond that!

Thanks!